



# **Chapter 3 – Part 2 Review (Introduction to SQL (2))**



# Null Values and Three Valued Logic

true, false, unknown

Any comparison with *null* returns *unknown*

Example:  $5 < \text{null}$  or  $\text{null} <> \text{null}$  or  $\text{null} = \text{null}$

Three-valued logic using the *unknown*:

OR:  $(\text{unknown} \text{ or } \text{true}) = \text{true}$ ,  
 $(\text{unknown} \text{ or } \text{false}) = \text{unknown}$   
 $(\text{unknown} \text{ or } \text{unknown}) = \text{unknown}$

AND:  $(\text{true} \text{ and } \text{unknown}) = \text{unknown}$ ,  
 $(\text{false} \text{ and } \text{unknown}) = \text{false}$ ,  
 $(\text{unknown} \text{ and } \text{unknown}) = \text{unknown}$

NOT:  $(\text{not } \text{unknown}) = \text{unknown}$

“*P* is unknown” evaluates to true if predicate *P* evaluates to *unknown*

Result of **where** clause predicate is treated as *false* if it evaluates to *unknown*



# Null Values and Three Valued Logic

## *Instructor*

| ID    | Name     | Dept_name | Salary |
|-------|----------|-----------|--------|
| 22222 | Einstein | Physics   | 95000  |
| 32343 | El Said  | History   |        |
| 33456 | Gold     | Physics   |        |

**SELECT name FROM instructor WHERE salary >100 and dept\_name = 'Physics';**

|              |                 |                |              |
|--------------|-----------------|----------------|--------------|
| <b>22222</b> | <b>Einstein</b> | <b>Physics</b> | <b>95000</b> |
|--------------|-----------------|----------------|--------------|

**SELECT name FROM instructor WHERE salary >100 or dept\_name = 'Physics';**

|              |                 |                |              |
|--------------|-----------------|----------------|--------------|
| <b>22222</b> | <b>Einstein</b> | <b>Physics</b> | <b>95000</b> |
| <b>33456</b> | <b>Gold</b>     | <b>Physics</b> |              |



# Aggregate Functions (Cont.)

Find the average salary of instructors in the Computer Science department

```
select avg (salary)  
from instructor  
where dept_name= 'Comp. Sci.';
```

Find the total number of instructors who teach a course in the Spring 2010 semester

```
select count (distinct ID)  
from teaches  
where semester = 'Spring' and year = 2010
```

Find the number of tuples in the *course* relation

```
select count (*)  
from course;
```



# Aggregate Functions – Group By

Find the average salary of instructors in each department

```
select dept_name, avg (salary)
from instructor
group by dept_name;
```

Note: departments with no instructor will not appear in result

| ID    | name       | dept_name  | salary |
|-------|------------|------------|--------|
| 76766 | Crick      | Biology    | 72000  |
| 45565 | Katz       | Comp. Sci. | 75000  |
| 10101 | Srinivasan | Comp. Sci. | 65000  |
| 83821 | Brandt     | Comp. Sci. | 92000  |
| 98345 | Kim        | Elec. Eng. | 80000  |
| 12121 | Wu         | Finance    | 90000  |
| 76543 | Singh      | Finance    | 80000  |
| 32343 | El Said    | History    | 60000  |
| 58583 | Califieri  | History    | 62000  |
| 15151 | Mozart     | Music      | 40000  |
| 33456 | Gold       | Physics    | 87000  |
| 22222 | Einstein   | Physics    | 95000  |

| dept_name  | avg_salary |
|------------|------------|
| Biology    | 72000      |
| Comp. Sci. | 77333      |
| Elec. Eng. | 80000      |
| Finance    | 85000      |
| History    | 61000      |
| Music      | 40000      |
| Physics    | 91000      |



# Aggregate Functions – Group By

고객별로 주문한 도서의 총 수량과 총 판매액을 구하시오

```
SELECT custID, COUNT(*) AS 도서수량, SUM(price) AS 총액
FROM orders
GROUP BY custID;
```

orders

| orderID | custID | bookID | price | order_date |
|---------|--------|--------|-------|------------|
| 2       | 1      | 3      | 21000 | 14/07/03   |
| 6       | 1      | 2      | 12000 | 14/07/07   |
| 1       | 1      | 1      | 6000  | 14/07/01   |
| 9       | 2      | 10     | 7000  | 14/07/09   |
| 3       | 2      | 5      | 8000  | 14/07/03   |
| 4       | 3      | 6      | 6000  | 14/07/04   |
| 10      | 3      | 8      | 13000 | 14/07/10   |
| 8       | 3      | 10     | 12000 | 14/07/08   |
| 7       | 4      | 8      | 13000 | 14/07/07   |
| 5       | 4      | 7      | 20000 | 14/07/05   |

| custID | 도서수량 | 총액    |
|--------|------|-------|
| 1      | 3    | 39000 |
| 2      | 2    | 15000 |
| 3      | 3    | 31000 |
| 4      | 2    | 33000 |



# Aggregate Functions – Having Clause

Find the names and average salaries of all departments whose average salary is greater than 42000

```
select dept_name, avg (salary)
from instructor
group by dept_name
having avg (salary) > 42000;
```

Note: predicates in the **having** clause are applied after the formation of groups whereas predicates in the **where** clause are applied before forming groups

```
SQL> select dept_name, count(*) from instructor where salary >1000 group by dept_name having count(*)>1;
```

| DEPT_NAME  | COUNT(*) |
|------------|----------|
| Physics    | 2        |
| Comp. Sci. | 3        |
| Finance    | 2        |
| History    | 2        |



## 3.8 Nested Subqueries

SQL provides a mechanism for the nesting of subqueries.

A **subquery** is a **select-from-where** expression that is nested within another query.

A common use of subqueries is to perform tests for set membership, set comparisons, and set cardinality.





# Nested Subqueries - Set Membership

Customer

| custID | Name |
|--------|------|
| 1      | 박지성  |
| 2      | 김연아  |
| 3      | 장미란  |
| 4      | 추신수  |
| 5      | 박세리  |

Orders

| orderID | custID |
|---------|--------|
| 2       | 1      |
| 6       | 2      |
| 1       | 1      |
| 9       | 3      |
| 3       | 2      |
| 4       | 3      |
| 10      | 4      |

**select** *name*  
**from** *customer*  
**where** *custID* **in** (**select** *custID* **from** *orders*);

IN에서 사용가능 한 subquery는  
결과로 다중 행, 다중열을 반환할 수 있음

\* Equivalent

→ **select** *name* **from** *customer*  
**where** *custID* = 1 **or** *custID* = 2 **or** *custID* = 3 **or** *custID* = 4;



# Nested Subqueries - Set Membership

Find courses offered in Fall 2009 and in Spring 2010

```
select distinct course_id
from section
where semester = 'Fall' and year= 2009 and
       course_id in (select course_id
                       from section
                       where semester = 'Spring' and year= 2010);
```

Find courses offered in Fall 2009 but not in Spring 2010

```
select distinct course_id
from section
where semester = 'Fall' and year= 2009 and
       course_id not in (select course_id
                              from section
                              where semester = 'Spring' and year= 2010);
```



# Nested Subqueries - Set Comparison

Find names of instructors with salary greater than that of some (at least one) instructor in the Biology department.

```
select distinct T.name  
from instructor as T, instructor as S  
where T.salary > S.salary and S.dept_name = 'Biology';
```

Same query using > **some** clause

```
select name  
from instructor  
where salary > some (select salary  
from instructor  
where dept_name = 'Biology');
```

| name | salary | Department |
|------|--------|------------|
| Lee  | 5000   | Biology    |
| Kim  | 10000  | Biology    |
| Park | 7000   | Computer   |

→ (5000, 10000)

결과는 ? Kim, Park



# Example Query of all Clause

Find the names of all instructors whose salary is greater than the salary of all instructors in the Biology department.

```
select name
from instructor
where salary > all (select salary
                        from instructor
                        where dept_name = 'Biology');
```

| name | salary | Department |
|------|--------|------------|
| Lee  | 5000   | Biology    |
| Kim  | 10000  | Biology    |
| Park | 7000   | Computer   |

결과는 ? No rows selected!!



# Exists

Customer

| custID | Name |
|--------|------|
| 1      | 박지성  |
| 2      | 김연아  |
| 3      | 장미란  |
| 4      | 추신수  |
| 5      | 박세리  |

Orders

| orderID | custID |
|---------|--------|
| 2       | 1      |
| 6       | 2      |
| 1       | 1      |
| 9       | 3      |
| 3       | 2      |
| 4       | 3      |
| 10      | 4      |

```
select name
from customer cs
where exists (select custID
                from orders od
                where cs.custID = od.custID);
```

| custID | Name |
|--------|------|
| 1      | 박지성  |
| 2      | 김연아  |
| 3      | 장미란  |
| 4      | 추신수  |



# Exists

## IN과 EXISTS의 차이

EXISTS : 단지 해당 row가 존재하는지만 확인하고, 더 이상 수행되지 않음 (True or False로 값을 Return)

IN : 실제 존재하는 데이터들의 모든 값까지 확인함  
(다중행 및 다중열로 값을 Return)

→ 일반적인 경우에 EXISTS가 더 좋은 성능 나타냄

## EXISTS 실행 순서

1. Customer cs 테이블에서 첫 행을 가지고 온 후 Subquery에 cs 값으로 입력
2. Subquery에서 Order od 테이블의 어떤 행에서 cs의 custID와 같은것을 찾으면 EXISTS는 True를 Return
3. 이후 cs테이블에서 첫 행에 대한 name이 반환

Customer cs 테이블에서 나머지 행에 대해서 반복



# Exists

Yet another way of specifying the query “Find all courses taught in both the Fall 2009 semester and in the Spring 2010 semester”

바깥질의 { **select** *course\_id*  
**from** *section* **as** S  
**where** *semester* = 'Fall' **and** *year* = 2009 **and**  
**exists** ( **select** \*  
하위질의 **from** *section* **as** T  
**where** *semester* = 'Spring' **and** *year* = 2010  
**and** S.*course\_id* = T.*course\_id*);

correlated

**Correlated subquery**

**Correlation name** or **correlation variable**



# Subqueries in the From Clause

SQL allows a subquery expression to be used in the **from** clause

Find the average instructors' salaries of those departments where the average salary is greater than \$42,000.

```
select dept_name, avg_salary
from ( select dept_name, avg (salary) as avg_salary
       from instructor
       group by dept_name)
where avg_salary > 42000;
```

select-from-where  
결과가 결국 또 다른  
relation이기 때문에  
가능한 것임

Note that we do not need to use the **having** clause → 바깥쪽에 where 절 때문

Another way to write above query

```
select dept_name, avg_salary
from ( select dept_name, avg (salary)
       from instructor
       group by dept_name)
       as dept_avg (dept_name, avg_salary)
where avg_salary > 42000;
```

from 절 안에 있는 select 구문의 결과를  
새로운 relation (즉, dept\_avg)로 재명명  
→ Oracle에서는 지원되지 않음





# With Clause

The **with** clause provides a way of defining a temporary view whose definition is available only to the query in which the **with** clause occurs.

Find all departments with the maximum budget

```
with max_budget (value) as
    (select max(budget)
     from department)
select budget
from department, max_budget
where department.budget = max_budget.value;
```

} with절을 포함한 질의에서만 유효한  
임시 relation (*max\_budget*) 생성



# Scalar Subquery

Scalar subquery is one which is used where a single value is expected

스칼라 값은 벡터 값에 대응되는 말로 단일값을 의미함!!

E.g. **select dept\_name,**      ↙ Subquery의 결과가 table이 아니라 value인 경우

```
(select count(*)  
  from instructor  
 where department.dept_name = instructor.dept_name)
```

```
  as num_instructors  
from department;
```

질의의 의미 ? → 학과별 교수님의 수를 출력



## Deletion (Cont.)

Delete all instructors whose salary is less than the average salary of instructors

```
delete from instructor  
where salary < (select avg (salary) from instructor);
```

Problem: as we delete tuples from instructor, the average salary changes

Solution used in SQL:

1. First, compute **avg** salary and find all tuples to delete
2. Next, delete all tuples found above (without recomputing **avg** or retesting the tuples)



## Insertion (Cont.)

Add all instructors to the *student* relation with *tot\_creds* set to 0

**insert into** *student*

```
select ID, name, dept_name, 0  
from instructor
```

insert 수행전에 select가  
먼저 수행되어야 함

The **select from where** statement is evaluated fully before any of its results are inserted into the relation (otherwise queries like

**insert into** *table1* **select \* from** *table1*

would cause problems, if *table1* did not have any primary key defined)

만약 **select from where** 구문을 수행하면서 **insert**문을 수행하면 어떻게 될까?  
즉, **select**를 해서 하나의 tuple을 뽑고 이를 다시 같은 **table**에 **insert**를 하고,  
주키 제약조건마저 없다면 어떻게 될까 ?

→ 중복된 tuple들이 계속해서 들어가고 결국 무한 loop에 빠진다.



# Modification of the Database – Updates

Increase salaries of instructors whose salary is over \$100,000 by 3%, and all others receive a 5% raise

Write two **update** statements:

① { **update instructor**  
**set salary = salary \* 1.03**  
**where salary > 100000;**

② { **update instructor**  
**set salary = salary \* 1.05**  
**where salary <= 100000;**

②→① 순서로 수행되면 ?

\$100000 보단 약간 적은 연봉  
(즉, 5% 인상하면 \$100000가  
넘는 연봉)을 받는 사람들은  
다시 3% 더 인상됨 → 본래 의도?

The order is important

Can be done better using the **case** statement (next slide)



# Case Statement for Conditional Updates

Same query as before but with case statement

**update** *instructor*

**set** *salary* = **case**

**when** *salary* <= 100000 **then** *salary* \* 1.05

**else** *salary* \* 1.03

**end**